

## 1. Introduction

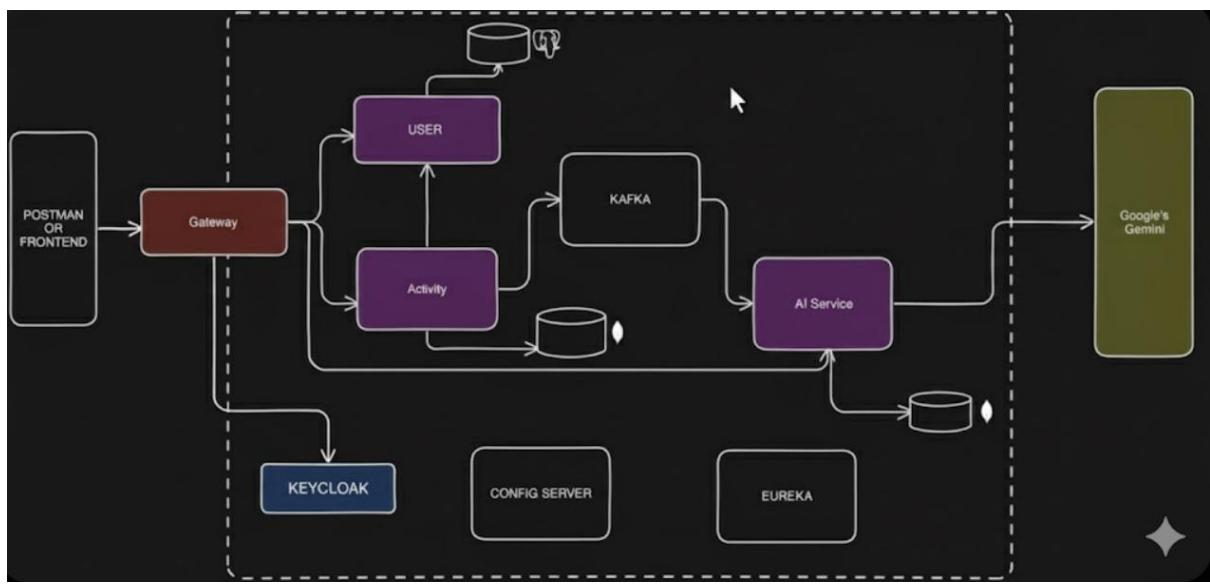
In today's world, fitness applications are becoming increasingly intelligent and data-driven.

Users expect personalized insights, real-time tracking, and AI-powered recommendations. To meet these expectations, our project implements a fully functional **AI-powered Fitness Recommendation System** using **Spring Boot Microservices, React, Kafka, and Google Gemini API**.

The backend of the system is designed using **six independent microservices**, each with a specific responsibility. This architecture improves scalability, security, independent deployment, and fault isolation. The system takes user activity logs (for example: walking, running, cycling), processes them through the AI microservice, and delivers personalized feedback, improvement areas, safety guidelines, and suggested workouts.

This report details the **design, implementation, technologies, and workflows** used in building the backend of this application.

## 2. System Architecture



The project follows a modern microservices architecture with the following core components:

### 2.1 List of All Backend Services

1. **User Service** – Manages user accounts, data storage, authentication integration.

2. **Activity Service** – Handles user fitness activities and publishes Kafka messages.
3. **AI Service** – Processes Kafka events and generates AI-driven recommendations.
4. **API Gateway** – Single entry point for all client requests, routing and token validation.
5. **Eureka Server** – Service discovery and registration.
6. **Config Server** – Centralized configuration storage for all services.

The system uses both **synchronous (REST)** and **asynchronous (Kafka)** communication.

User → API Gateway → Microservice calls are synchronous. Activity → Kafka → AI Service flow is asynchronous for scalability and performance.

## 2.2 Architecture Features

- **Loose coupling:** Each service runs independently on different ports.
- **Decentralized data management:**
  - PostgreSQL for User Service
  - MongoDB for Activity & AI Service
- **Event-driven processing** using Kafka.
- **Security** implemented using **Keycloak OAuth2 + JWT**.
- **Automatic service registration** handled by Eureka.
- **Central configuration** managed by Config Server.

This architecture reflects real industry standards used in banking, fitness, e-commerce, and enterprise applications.

## 3. Technology Stack

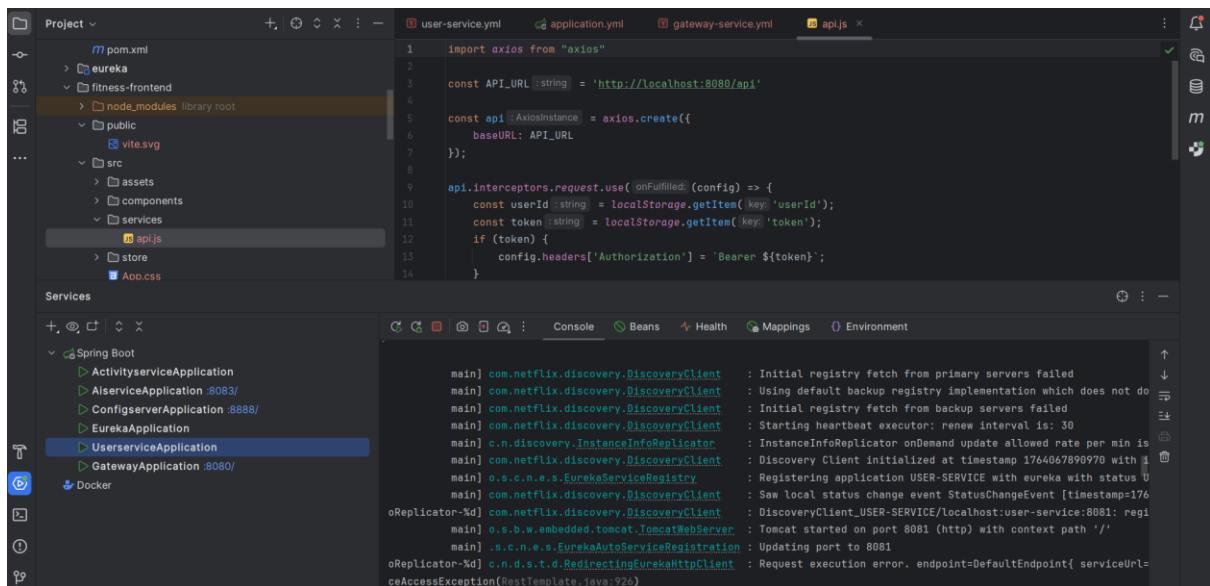
### 3.1 Backend Technologies

Component	Technology
Microservices	Spring Boot
Routing	Spring Cloud Gateway

Authentication	Keycloak OAuth2 / OpenID Connect
Service Discovery	Eureka (Spring Cloud Netflix)
Asynchronous Messaging	Apache Kafka
Databases	PostgreSQL, MongoDB
Build Tool	Maven
AI Integration	Google Gemini API

### 3.2 Reason for Technology Choices

- **Spring Boot** simplifies microservice creation with built-in support for REST, JPA, and cloud tools.
- **PostgreSQL** is robust for relational data like users.
- **MongoDB** is ideal for flexible activity logs and AI responses.
- **Kafka** enables high-speed asynchronous processing.
- **Keycloak** provides enterprise-level authentication and authorization.
- **Gemini API** allows free AI model usage for generating recommendations.



The screenshot shows a Java IDE interface with the following details:

- Project View:** Shows the project structure with files like `pom.xml`, `eureka`, `fitness-frontend`, `node_modules`, `public`, `vite.svg`, `src`, `assets`, `components`, `services` (containing `api.js`), `store`, and `App.css`.
- Code Editor:** Displays a file named `apis` containing the following code:

```

1 import axios from "axios"
2
3 const API_URL :string = 'http://localhost:8080/api'
4
5 const api :AxiosInstance = axios.create({
6   baseURL: API_URL
7 });
8
9 api.interceptors.request.use( onFulfilled: (config) => {
10   const userId :string = localStorage.getItem( key: 'userId' );
11   const token :string = localStorage.getItem( key: 'token' );
12   if (token) {
13     config.headers['Authorization'] = `Bearer ${token}`;
14   }
}

```
- Services View:** Shows a list of services including `ActivitiServiceApplication`, `AlserviceApplication :8083/`, `ConfigserverApplication :8888/`, `EurekaApplication`, **UserserviceApplication** (selected), and `GatewayApplication :8080/`.
- Console View:** Shows the application's log output, including Eureka client registration messages and Tomcat startup logs.

## 4. Detailed Implementation

This section describes the backend implementation of each component in detail.

### 4.1 Config Server

The screenshot shows a code editor with two panes. The left pane displays a project structure for a Spring Boot application. It includes directories for activityservice, aiservice, configserver, eureka, fitness-frontend, gateway, and userservice. Under configserver, there's a .mvn directory, a src directory containing main (with java and resources sub-directories), and a test directory. The resources/main/config folder contains several YAML files: activity-service.yml, ai-service.yml, gateway-service.yml, user-service.yml, and application.yml. The right pane shows a snippet of the application.yml file:

```

3   name: configserver
4   profiles:
5     active: native
6   cloud:
7     config:
8       server:
9         native:
10        search-location: classpath:/config
11
12
13  server:
14    port: 8888

```

The Config Server centralizes all configuration files (application.yml).

All microservices fetch their configurations at startup.

This avoids duplication and improves maintainability.

The Config Server was created using Spring Cloud Config starter and linked to a configuration repository.

## 4.2 Eureka Server

The screenshot shows a web browser window displaying the Eureka dashboard. The URL is localhost:8761. The page has a green header bar with the text "EMERGENCY! EUREKA MAY BE INCORRECTLY CLAIMING INSTANCES ARE UP WHEN THEY'RE NOT. RENEWALS ARE LESSER THAN THRESHOLD AND HENCE THE INSTANCES ARE NOT BEING EXPIRED JUST TO BE SAFE." Below this, there are sections for "DS Replicas" and "General Info".

**DS Replicas**

Instances currently registered with Eureka

Application	AMIs	Availability Zones	Status
ACTIVITY-SERVICE	n/a (1)	(1)	UP (1) - <a href="#">localhost:activity-service:8082</a>
AI-SERVICE	n/a (1)	(1)	UP (1) - <a href="#">localhost:ai-service:8083</a>
GATEWAY-SERVICE	n/a (1)	(1)	UP (1) - <a href="#">localhost:gateway-service:8080</a>
USER-SERVICE	n/a (1)	(1)	UP (1) - <a href="#">localhost:user-service:8081</a>

**General Info**

Name	Value
total-avail-memory	104mb
num-of-cpus	16
current-memory-usage	43mb (41%)
server-uptime	03:34
registered-replicas	

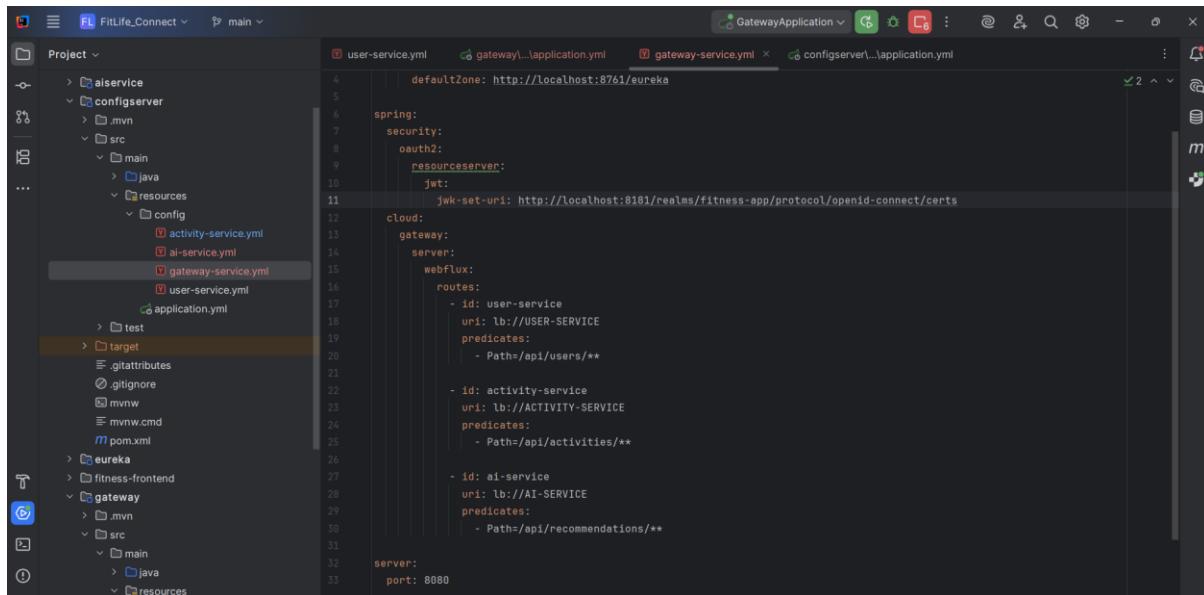
The Eureka Server acts as the service registry.

Each microservice automatically registers itself using its service name.

The Gateway queries Eureka to discover where each service is running.

This removes hard-coded URLs and allows dynamic routing even if services restart.

### 4.3 API Gateway



```
defaultZone: http://localhost:8761/eureka
spring:
  security:
    oauth2:
      resourceServer:
        jwt:
          jwk-set-uri: http://localhost:8181/realm/fitness-app/protocol/openid-connect/certs
cloud:
  gateway:
    server:
      webflux:
        routes:
          - id: user-service
            uri: lb://USER-SERVICE
            predicates:
              - Path=/api/users/**
          - id: activity-service
            uri: lb://ACTIVITY-SERVICE
            predicates:
              - Path=/api/activities/**
          - id: ai-service
            uri: lb://AI-SERVICE
            predicates:
              - Path=/api/recommendations/**
server:
  port: 8080
```

The API Gateway acts as the secure edge layer of the system.

Features implemented:

- All client calls pass through the gateway.
- Keycloak token validation using OAuth2.
- Routing to internal services based on URL patterns.
- Filters for security and request logging.

Routing example:

- /api/user/\*\* → User Service
- /api/activity/\*\* → Activity Service
- /api/ai/\*\* → AI Service

This architecture hides internal microservices from direct exposure.

### 4.4 User Service (PostgreSQL)

The User Service stores and manages user profiles.

Technologies used:

- Spring Boot Web
- Spring Data JPA
- PostgreSQL
- Lombok
- Validation annotations

#### **Key Features:**

- UUID-based user ID generation
- Password & email validation
- Automatic timestamps using @CreationTimestamp
- Data stored using PostgreSQL table users

The service also integrates with Keycloak for secure login/logout.

#### **4.5 Activity Service (MongoDB)**

The Activity Service handles all user activity logs.

#### **Functions implemented:**

- Add new activity
- Retrieve all activities for a user
- Publish activity events to Kafka

MongoDB was used because activity logs may vary in format and grow continuously.

#### **Kafka Publishing**

Whenever a user logs an activity, a message is sent to Kafka topic:

activity-topic

The message contains:

- userId

- activity type
- calories burned
- duration
- timestamp

This triggers asynchronous processing in the AI Service.

## 4.6 Kafka Integration

Kafka is deployed using Docker.

Two components are used:

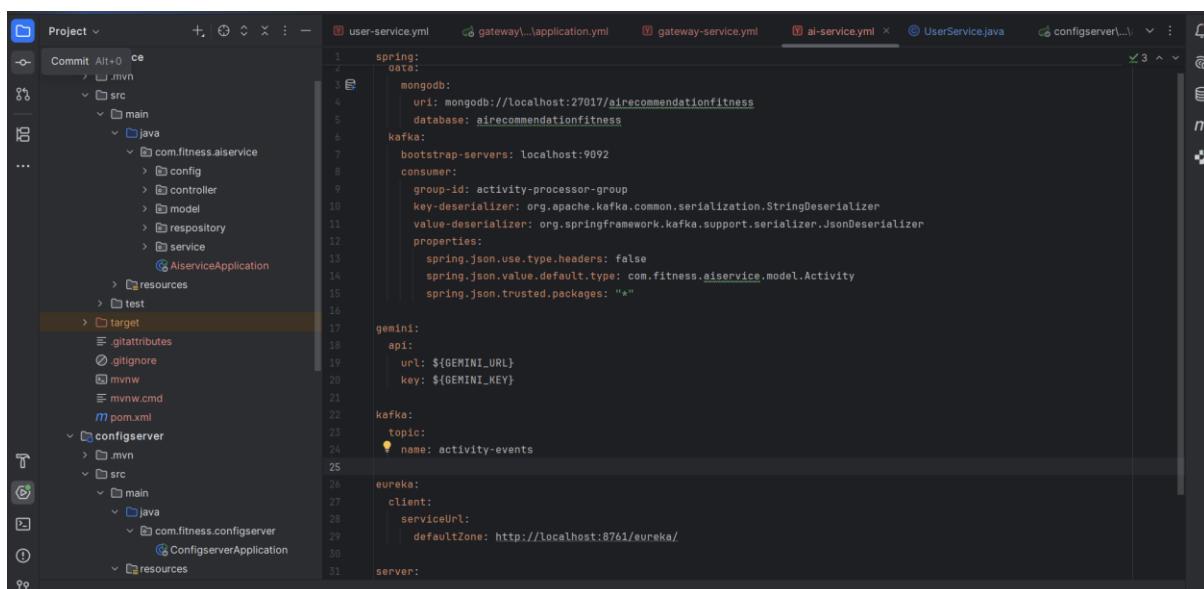
- Kafka Broker
- Zookeeper

### Why Kafka?

- High throughput
- Asynchronous event processing
- Services remain loosely coupled
- AI Service can process activities without delaying the user

This design ensures the system stays responsive and scalable.

## 4.7 AI Service (Gemini API + MongoDB)



```

spring:
  data:
    mongo:
      url: mongodb://localhost:27017/airecommendationfitness
      database: airecommendationfitness
    kafka:
      bootstrap-servers: localhost:9092
      consumer:
        group-id: activity-processor-group
        key-deserializer: org.apache.kafka.common.serialization.StringDeserializer
        value-deserializer: org.springframework.kafka.support.serializer.JsonDeserializer
        properties:
          spring.json.use.type.headers: false
          spring.json.value.default.type: com.fitness.aiservice.model.Activity
          spring.json.trusted.packages: "*"
    gemini:
      api:
        url: ${GEMINI_URL}
        key: ${GEMINI_KEY}
      kafka:
        topic:
          name: activity-events
    eureka:
      client:
        serviceUrl:
          defaultZone: http://localhost:8761/eureka/
      server:

```

The AI Service is the intelligent component of the system.

### **Workflow:**

1. AI Service consumes activity events from Kafka.
2. Extracts required activity information.
3. Calls **Google Gemini AI API**.
4. Receives AI-generated recommendations.
5. Saves data to MongoDB.
6. Sends response to frontend via REST API.

### **Gemini Output Includes:**

- Overall analysis
- Improvement suggestions
- Safety guidelines
- Activity summary

The AI Service transforms raw activity logs into personalized fitness insights.

## **5. Results & Conclusion**

The backend system is fully developed with all six services functioning independently.

The microservices interact using modern enterprise-level patterns such as service discovery, API gateway routing, JWT-based authentication, and Kafka event processing.

The AI Service successfully generates personalized recommendations using Google Gemini API.

### **5.1 Achievements**

- Built a scalable, maintainable microservices architecture.
- Implemented event-driven processing using Kafka.
- Integrated AI using a free and powerful model.
- Designed two different databases for optimal data handling.

- Created a secure ecosystem with Keycloak authentication.

## 5.2 Future Enhancements

- Deployment using Docker Compose or Kubernetes.
- Adding caching using Redis.
- Enhancing AI recommendations with custom model training.
- Creating an analytics dashboard in the UI.